

Compiler Testing via a Theory of Sound Optimisations in the C11/C++11 Memory Model

Robin M., Pankaj P., Francesco Z.

Presented by
Akshay Gopalakrishnan

December 19, 2021

Introduction

Testing sequential programs to hunt for compiler bugs is well established at this time. However, testing concurrency bugs still remains a hard problem. This is majorly due to ill understood specification of the memory consistency model, and lack of information on the various transformations that preserve concurrent program semantics. This work designs a strategy to reduce the complexity of finding concurrency bugs in C11/C++11 via testing sequential code. The source code trace is compared with that of the final end code trace. Previous theoretical work establishing soundness of various local transformations in concurrent execution is utilized to achieve this.

Example of Compiler optimisation

Consider the following example C code, where *th_1* and *th_2* are run by two different threads.

```
int g1 = 1; int g2 = 0;

void *th_1(void *p) {
    for (int l = 0; (l != 4); l++) {
        if (g1) return NULL;
        for (g2 = 0; (g2 >= 26); ++g2)
            ;
    }
}

void *th_2(void *p) {
    g2 = 42;
    printf("%d\n", g2);
}
```

The above code should print the value of *g2* to be 42.

However, the above code when compiled using gcc 4.7.0 with -O2 enabled on x86-64 machine, the result of printf is 0.

Compiler Testing via a Theory of Sound Optimisations in the C11/C++11 Memory Model

Example of Compiler optimisation

Consider the following example C code, where `th_1` and `th_2` are run by two different threads.

```
int g1 = 0, g2 = 0;  
void th_1(void*) {  
    int i; for (i = 0; i < 100; i++)  
        g1 = g1 + 1; g2 = g2 + 1;  
}
```

The above code should print the value of `g2` to be 42. However, the above code when compiled using gcc 4.7.0 with `-O2` enabled on x86-64 machine, the result of `printf` is 0.

The reason why it should be 42 is that the value of `g2` is never changed in `th_1`. Notice that `g1` has 1, so the conditional will always return *NULL*. It is only `th_2` that writes to `g2`.

The above code is optimized using Loop invariant code motion in *th_1*. However, the code that results due to it in x86-64 format is as below:

```
th_1:
    movl  g1(%rip), %edx    # load g1 (1) into edx
    movl  g2(%rip), %eax    # load g2 (0) into eax
    testl %edx, %edx       # if g1 != 0
    jne   .L2              # jump to .L2
    movl  $0, g2(%rip)
    ret
.L2:
    movl  %eax, g2(%rip)    # store eax (0) into g2
    xorl  %eax, %eax        # return 0 (NULL)
    ret
```

Notice the last *movl* before return. Here the value of *g2* is restored to be 0. This is sound if the program was sequential, however, in a concurrent setting this is clearly unsafe.

Compiler Testing via a Theory of Sound Optimisations in the C11/C++11 Memory Model

└ Behind the scenes

The above code is optimized using Loop invariant code motion in `th_1`. However, the code that results due to it in x86-64 format is as below:

```

w1:1
movl    %0(%rip), %edx    # load %0 into edx
xorl    %0(%rip), %edx    # load %0 into edx and
xorl    %edx, %edx         # set %edx to 0
jmp     .L2               # jump to .L2

.L2:
movl    %eax, %0(%rip)    # store eax to %0
xorl    %eax, %eax        # return 0 (NULL)
ret

```

Notice the last `movl` before `return`. Here the value of `g2` is restored to be 0. This is sound if the program was sequential, however, in a concurrent setting this is clearly unsafe.

Here, the optimization tries to eliminate the inner loop. However, the process introduces some redundant writes, like that of $g2$ to be present. Thus, redundant write introduction in this case becomes unsafe. This in turn makes such a variant of loop invariant code motion unsafe.

Main idea of testing using traces: Another Example

Testing for such concurrency bugs is difficult. For this purpose, we note that the concurrency bugs are only among those actions which involve a shared memory. So reducing the source program to a set of actions performed on shared memory can first be done. This is followed by identifying the possible traces of these actions allowed by the program on its executions. The above code's problem can be observed by the following trace observed to be incorrect after optimization.

Init	g1	1
Init	g2	0
Load	g1	1

Init	g1	1
Init	g2	0
Load	g1	1
Load	g2	0
Store	g2	0

Compiler Testing via a Theory of Sound Optimisations in the C11/C++11 Memory Model

└ Main idea of testing using traces: Another Example

The memory model of C11 by this time was defined using traces / execution based semantics. It was more of utilizing this to do practical testing of C programs.

Testing for such concurrency bugs is difficult. For this purpose, we note that the concurrency bugs are only among those actions which involve a shared memory. So reducing the source program to a set of actions performed on shared memory can first be done. This is followed by identifying the possible traces of these actions allowed by the program on its executions. The above code's problem can be observed by the following trace observed to be incorrect after optimization.

Proc	P1		P2
Exec	g1		g2
Exec	g1		g2
Exec	g1		g2
Exec	g1		g2

Overview of C11 memory model

To summarize the memory model, each shared memory access has a memory order. Each action has an access, memory order and thread id.

$$\text{mem_ord}, \mu ::= \text{NA} \mid \text{SC} \mid \text{ACQ} \mid \text{REL} \mid \text{R/A} \mid \text{RLX}$$
$$\phi ::= \text{R}_\mu l v \mid \text{W}_\mu l v \mid \text{Lock } l \mid \text{Unlock } l \mid \text{Fence}_\mu \mid \text{RMW}_\mu l v_1 v_2$$
$$\text{actions} ::= aid, tid: \phi$$

Useful terminology used for optimization

- Every concurrent program has various executions possible.
- These executions involve various shared memory *actions* that are done (even those due to multiple iterations of loop).
- A collection of these actions is defined to be an *opsem*.
- *opsem* also retains the syntactic order between actions belonging to same thread.
- All possible *opsems* for a program is called *opsemset*.

An example is as below;

`x = 1; y = 1; if (x == y){x = 42;}`
has, among others, the two opsems below:



Effect of an optimization

- The effect of an optimization is directly seen at the *opsem* level.
- For an *opsem*, an optimization has the effect of reordering, eliminating and introducing actions.

The following example showcases this:

```
for (i=0; i<2; i++) {  
    z = z + y + i;  
    x = y;  
}  
⇒  
t = y; x = t;  
for (i=0; i<2; i++) {  
    z = z + t + i;  
}
```

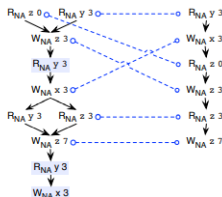


Figure 3: effect of loop invariant code motion (LIM) on an opsem

Soundness of Optimisations

- Every *opsem* represents an execution.
- The outcome of the execution can be characterized as the final values in memory and the final read values.
- These two are represented together as an *obs* or observable behavior.
- A pair *opsem*, *obs* represent a candidate execution.
- An optimization is sound if the set of *obs* of the transformed program is a subset of that of *obs*.

Types of Optimizations considered

- Elimination - eg: Read-after-Write, Write-after-Write, etc.
- Reordering - eg: Read-Write, Read-Read, etc.
- Introduction - Read / Write.

Using above results for testing

Results

Thank you

Questions?