

Explaining Relaxed Memory Models with Program Transformations

Ori Lahav and Viktor Vafeiadis

Presented by
Akshay Gopalakrishnan

September 26, 2021

Introduction

- Uniprocessor semantics respects sequential consistency.
- Multiprocessor has more behaviors that programs can exhibit.
- They are either justified by hardware features such as Load/Store buffers, speculation, etc.
- Or they can be justified by the program being optimized in various passes by a compiler.
- Semantics of such behaviors however, are not specified using such intuition, rather, per-execution or event-structure based axiomatic models, or more non-trivial operational models.
- This paper investigates upto what extent such memory consistency models can be specified using program transformations.

- TSO can be explained by W-R reordering and Read-after-Write elimination over SC.
- RA cannot be specified in the same manner - counter example to show this.
- A substantial subset of POWER can be specified in a similar way over a stronger model of power - SPOWER
- ARM however, cannot be done the same way.
- Advantage of using this to simplify compilation correctness proofs.

Axiomatic Model Definitions

Axiomatic events:

- Read event $\mathcal{R} = Ld \cup U$.
- Write event $W = St \cup U$.
- Fence - Can be of various types depending on the architecture / language's memory model.

Binary relations:

- **po** - Program order per thread.
- **rf** - Between a read and the write from which its read value comes.
- **co** - Order between writes/updates to same location.
- **mo** - Order between memory writes/updates.

Sequential Consistency (SC)



Additional binary relations:

- **hb** - Happens before $hb = (po \cup rf)^+$
- **fr** - Reads before (from reads) $fr = (rf^{-1}; mo|_{loc})$

SC is defined using the following irreflexivity relations that must hold for any program execution:

- | | |
|----------------------------------|---------------------------------|
| a) mo total order on W/U. | b) hb irreflexive. |
| c) mo;hb irreflexive. | d) fr;hb irreflexive. |
| e) fr;mo irreflexive. | f) fr;mo;hb irreflexive. |

Figures to explain above axioms of irreflexivity



Total Store Order (TSO)



Additional binary relations:

- xhb - Happens before $xhb = (po \cup rf)^+$
- fr - Reads before (from reads) $fr = (rf^{-1}; mo|_{loc})$
- rfe - Reads from external $rf \setminus po$

x86 TSO is defined using the following irreflexivity relations that must hold for any program execution:

- a) mo is a strict total order on W/U/F.
- b) xhb irreflexive.
- c) $mo; xhb$ irreflexive.
- d) $fr; xhb$ irreflexive.
- e) $fr; mo$ irreflexive.
- f) $fr; mo; rfe; po$ irreflexive.
- g) $fr; mo; [U \cup F]; po$ irreflexive.

Figures to explain the last two axioms of irreflexivity.



TSO Equivalence

TSO is precisely categorized by write-read reordering and read-after-write elimination over SC.

Theorem 1

*A plain execution G is TSO-consistent iff $G \mapsto^*_{TSO} G'$ for some SC-consistent execution G' .*

where \mapsto^*_{TSO} implies we either do write-read reordering or read-after-write elimination.

Proposition 1

If $G \mapsto_{TSO}^ G'$ and G' is TSO consistent, then so is G .*

Proposition 2

If G is TSO-consistent execution, then so is $\text{RemoveWR}(G, a, b)$.

Proposition 3

If G is TSO-consistent execution, then so is $\text{ReorderWR}(G, a, b)$ if $\langle a, b \rangle \notin \text{mo}; \text{rf}$.

Proposition 4

Suppose G is TSO-consistent but not SC-consistent, then $G \mapsto_{TSO}^ G'$ for some TSO-consistent execution G' .*

Release Acquire (RA)

This model exhibits non-multi-copy-atomicity, meaning, threads may observe updates/writes in different orders. Here **mo** is defined only between events W/U accessing the same memory. The axioms then are as follows:

- a) **mo** disjoint union of relations mo_x each of which are strict total orders.
- b) **hb** irreflexive.
- c) **mo;hb** irreflexive.
- d) **fr;hb** irreflexive.
- e) **fr;mo** irreflexive.

where **hb**, **fr** are as that of Sequential consistency.

Why RA cannot be described the same way

$$\begin{array}{l} a := x; \textcolor{teal}{// 1} \\ b := y; \textcolor{teal}{// 0} \end{array} \parallel x := 1; \parallel y := 1; \parallel \begin{array}{l} c := y; \textcolor{teal}{// 1} \\ d := x; \textcolor{teal}{// 0} \end{array}$$

$$\begin{array}{l} x := 1; \\ a := x; \textcolor{teal}{// 1} \\ b := y; \textcolor{teal}{// 0} \end{array} \parallel y := 1; \parallel \begin{array}{l} c := y; \textcolor{teal}{// 1} \\ d := x; \textcolor{teal}{// 0} \end{array} \rightsquigarrow \begin{array}{l} b := y; \textcolor{teal}{// 0} \\ x := 1; \\ a := 1; \textcolor{teal}{// 1} \end{array} \parallel y := 1; \parallel \begin{array}{l} c := y; \textcolor{teal}{// 1} \\ d := x; \textcolor{teal}{// 0} \end{array}$$

Why RA cannot be described the same way

$$\begin{array}{l} y := 1; \\ x := 1; \\ a := x; \text{ // } 3 \\ b := z; \text{ // } 0 \end{array} \parallel x := 3; \parallel \begin{array}{l} z := 1; \\ x := 2; \\ c := x; \text{ // } 3 \\ d := y; \text{ // } 0 \end{array}$$

Similar Analysis for POWER and ARMv7/v8

- POWER can be specified using SPOWER, whose only difference is it disallows $po \cup rf$ cycles.
- Then POWER is equivalent to reordering over SPOWER (you keep reordering until only dependency order = program order).
- ARM cannot be specified in a similar way (perhaps refer to figure in notes section that cannot be explained by program transformations)

Advantage of Transformational Models: Compiler Correctness

Consider $[[P]]_M$ to be behaviors of program P under memory model M . Assume we have a compilation scheme from source C to target A (code gen phase: direct mapping of instructions). Let M_C and M_A be their respective memory models. Then compiler correctness requires:

$$\forall P_C. [[compile(P_C)]]_{M_A} \subseteq [[P_C]]_{M_C}.$$

Compiler Correctness simplified using results in paper

Theorem 1 and 2 provide us with the following property:

$$\forall P_C. [[\text{compile}(P_C)]]_{M_A} \subseteq \cup \{ [[P'_A]]_{SM_A} \mid P'_A.\text{compile}(P_C) \mapsto^*_{M_A} P'_A \}.$$

where SM_A is the stronger memory model than M_A (like SC for TSO in Theorem 1).

Simplified proof requirements

Using above info, compiler correctness boils down to two conditions.

- Compilation should be correct for the strong model SM_A .
- There should be a set of source program transformations sound for source model M_C and it should capture all target transformations from a compiled program.

Thank you

Questions ?