

Laws of Order: Expensive Synchronization in Concurrent Algorithms Cannot be Eliminated

hagit A, Rachid G, Danny H, Petr K, Maged M, Martin V

Presented by
Akshay Gopalakrishnan

November 9, 2021

Introduction

- Designing Concurrent Algorithms with good performance is a non-trivial tasks.
- Major part of slowdown is due to synchronization.
- Hence efforts are made to remove such costly synchronizations.
- This paper shows that it is impossible to remove completely all expensive synchronization primitives for a class of concurrent implementations.
- Such a result helps designers to assert when they can stop improving an algorithm this way.

Summary of Results

The results are based on the usage of atomic Read-after-Write and Write-after-Read operations. The authors show that without using the above primitives, it is:

- Impossible to build a linearizable implementation that is non-commutative and respects deterministic sequential specification.
- Impossible to build an algorithm that respects mutual exclusion and is deadlock-free.

Mutual Exclusion

Given N processes, each of which accesses a critical section by acquiring a lock, mutual exclusion requires

No more than one process can be in the critical section at the same time

Proof intuition

Proof by contradiction. Part 1:

- Assume that we do not use a shared write updating the lock.
- Then this would imply more than one thread can be in the critical section.

Part 2:

- Assume that we do not use ARAW or AWAR as part of the lock implementation.
- Then this would imply more than one thread can read concurrently the same lock and be ready to acquire the lock (update the lock variable).
- Then one thread can acquire the lock (by updating the lock variable).
- After which another thread can also acquire the lock (by updating lock variable too!).
- Thu, violating mutual exclusion.

An algorithm is linearizable w.r.t a sequential specification if

Each execution of an algorithm is equivalent to some sequential execution of that specification, where the order between non-overlapping methods is preserved.

Note: Not all linearizable algorithms need to use RAW/WAR. Only some, which have two properties.

Two Properties of Linearizable Algorithms that must use RAW/WAR

- Deterministic Sequential Specification - Executing one step of program from same state should yield same result all the time.
- Strongly non-commutative methods - Methods whose execution influences further call to the same method.

Proof Intuition

For this, let us consider a simple example of a set data-structure, which has three methods:

- $contains(k) - \{ret = (k \in A) \wedge (S = A)\}.$
- $remove(k) - \{ret = (k \in A) \wedge (S = A \setminus \{k\})\}.$
- $add(k) - \{ret = (k \notin A) \wedge (S = A \cup \{k\})\}$

Firstly note that these methods respect Deterministic Sequential specification. Secondly, note that *add* and *remove* are strongly non-commutative. For instance, if we perform *add(k)* twice, then only one of them will return true, whichever is executed first. Similar reasoning for *remove*.

Proof intuition

Proof by contradiction:

- Suppose, such a linearizable algorithm does not use RAW or AWAR.
- Given that the methods are strongly non-commutative, they must perform some form of shared write to actually influence a subsequent call to the method.
- And also for them to be influenced by previous method calls, they must also perform a shared read.
- Now suppose $add(k)$ is performed, but not using AWAR or RAW.
- This means, we can have an execution where for instance, the method reads from the set, decides it can write but does not do it.
- Now another $add(k)$ does the same thing, but it does its operation fully. Thus writing to the set.
- Now the previous call to $add(k)$ can be continued to write another k to the set, which is incorrect.

Examples of strongly non-commutative methods

- *add* and *remove* methods in the above set example.
- *push* and *pop* method for stack / queue data structures.
- CAS - Compare and Swap.

Formal proof elements

- Formal language with transition semantics.
- Formal definition for executions.
- Notion of histories as sub-executions (traces).
- Proving both above claims in the same flow using these formal elements.

Conclusion

- Utilization of concurrent computation certainly does not come without its downsides.
- The use of RAW / AWAR cannot be avoided to ensure properties such as Mutual Exclusion and variants of Linearizability.

Thank you

Questions?