

# Understanding POWER Multiprocessors

Sushmit Sarkar, Peter Sewell, Jade Alglave, Luc Maranget,  
Derek Williams

Presented by  
Akshay Gopalakrishnan

November 10, 2021

# Introduction

- Much of the performance in hardware comes due to features such as Read/Write buffers, Speculation, Caches, etc.
- The behavior of execution of programs utilizing all such features in a hardware can be defined by a relaxed memory consistency model.
- ARM, x86, POWER.. all these hardware exhibit relaxed behaviors.
- Of these POWER is relatively well understood.
- Major reason is due to informal specification and behaviors described only via litmus tests.
- Continual hardware improvements in POWER also have resulted in more relaxed behaviors.

- This paper analyzes the POWER multiprocessor family for relaxed behaviors explain their discovered behaviors via an Abstract Machine.
- This helps in avoiding getting involved in the complexities of hardware itself while all the way understanding the weak behaviors exhibited by them.
- The abstract machine has only one storage subsystem while having read/write buffers for each thread.

# Preliminary Definitions


## Axiomatic events

- Read -  $R[x]$
- Write -  $W[x]$

## Binary Relations

- Program order - *po* (per thread syntactic order)
- Reads from - *rf* (from a write to a read whose value is the write)


# Example 1: Message Passing



Test MP : Allowed

The outcome in program is not allowed under sequential execution, but in POWER such an outcome is allowed. In real program, the read to y can be in a loop, which only ends once it reads 1, indicating that the value of x has been modified. But this behavior should then clearly not be possible. But it is.

## Example 2: Store Buffering




Test SB : Allowed

Similarly the outcome in the program above is also allowed under POWER. Such an outcome is also observable in x86-TSO.

## Example 3: IRIW


POWER also allows writes to be non-multi-copy-atomic (nmca). Meaning writes can be visible to processors in different orders. The following example showcases an outcome that should not be possible in multi-copy-atomic hardware.



Test IRIW : Allowed

## Example 4: WRC

Another flavor of nmca is the following example



Test WRC : Allowed

In this example, the read value of `x` is 1, meaning the write has been propagated to Thread 1. Subsequently, Thread 2 reads from a write `po` after the read to `x` in Thread 1. However, this does not ensure that Thread 2 has also received the updated value of `x`. The outcome above is allowed in POWER.



# Use of Synchronization and other Barriers

The above behaviors of programs may not be desirable in all situations. To ensure this, POWER also comes with barrier instructions: mainly *sync* and *lwsync*. Let us see how these barriers can be used to ensure ordering constraints.

